


# Eden: Developer-friendly User-space Far Memory

Anil Yelam, Stewart Grant, Saarth Deshpande, Nadav Amit, Radhika Niranjana Mysore, Amy Ousterhout, Marcos K. Aguilera and Alex C. Snoeren  
[ayelam@ucsd.edu](mailto:ayelam@ucsd.edu)



## Far memory systems

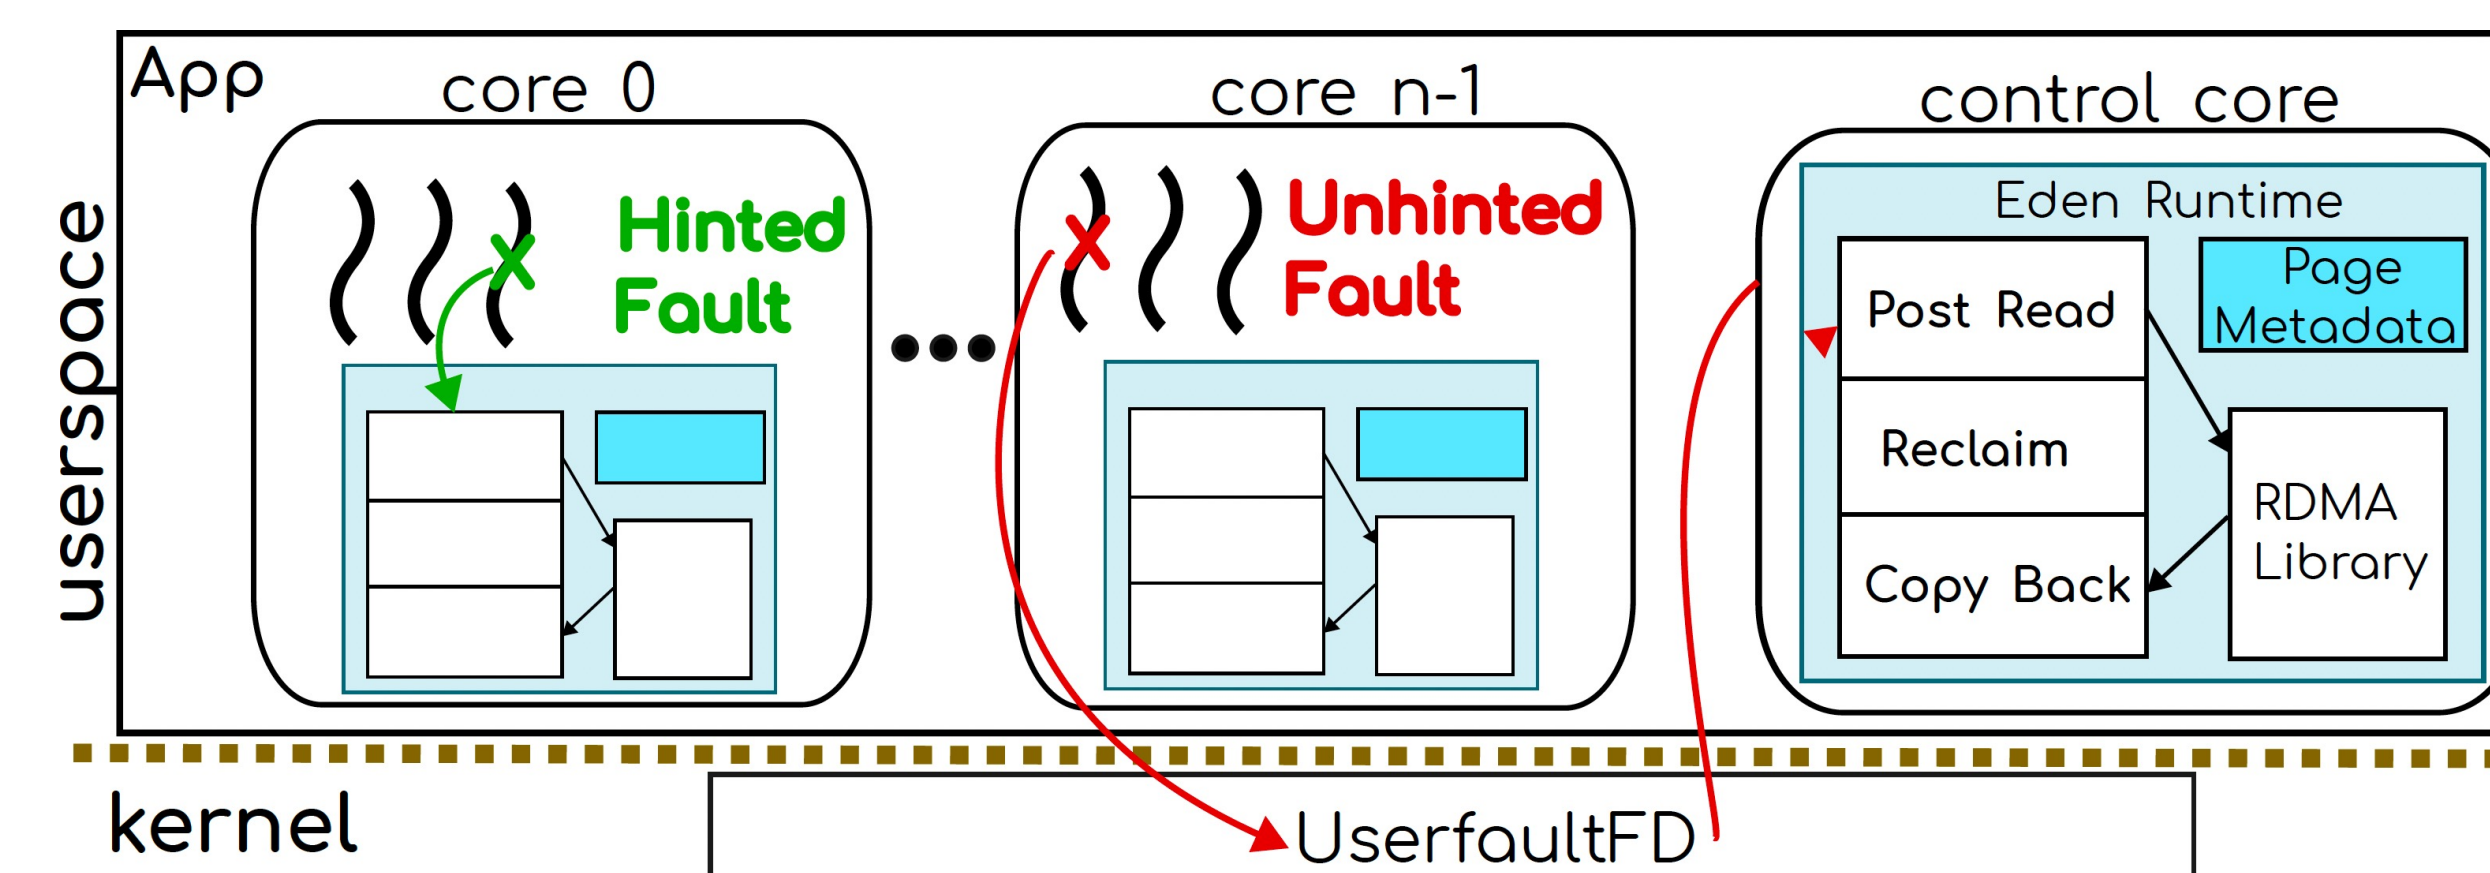
extend available memory on servers with (remote) memory from other servers.

## Current approaches

- **Kernel-based systems** build on OS swapping to remote memory. They suffer from the *paging overheads and rigid policies* e.g., **Fastswap** [Eurosys'20]
- **Runtime-based systems** enable flexible, app-specific memory management policies but require *significant application changes* e.g., **AIFM** [OSDI'21]

## Eden charts a middle-ground!


- Keeps the page-based virtual memory interface
- Manages pages from Userspace for flexibility
- Expects just a few "hints" for page faults



Eden's user-space runtime based on `userfaultfd` and Shenango [NSDI'19]

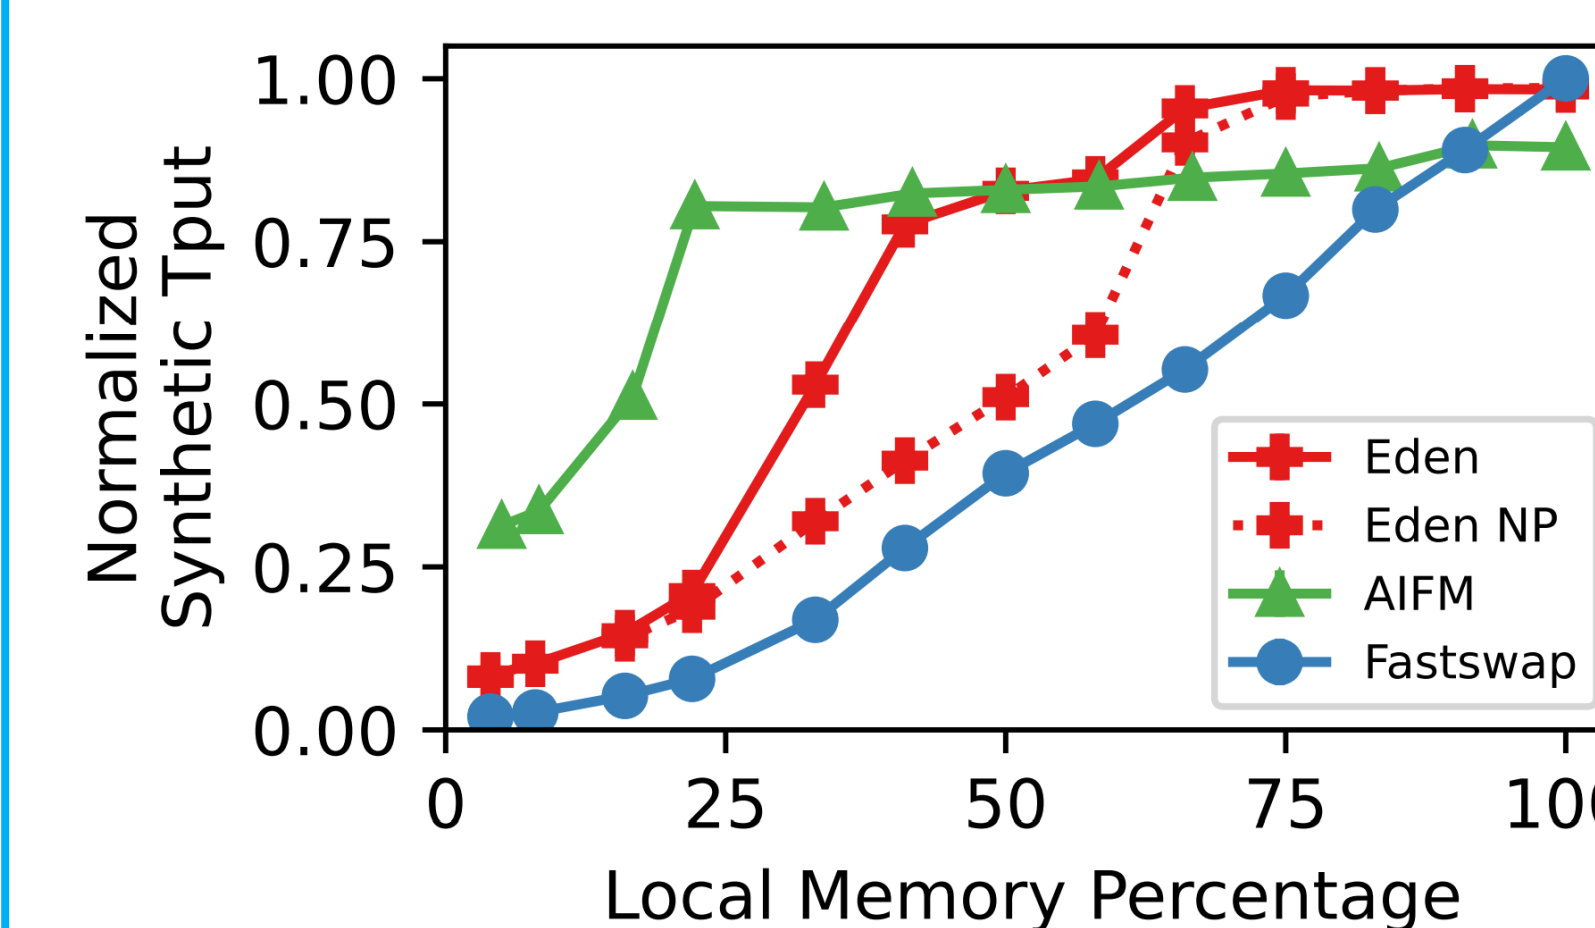
## Key result 1: Only a handful of hints needed!

For most apps, a small number of locations cover 95% of faults, even after the library code is excluded. Further results show that the set of locations is robust across different degrees of memory pressure.




## Key result 2: Comparable performance to AIFM but with significantly (50x) less effort

For two benchmarks taken from **AIFM**, Eden performs significantly better than **Fastswap** and stays close to **AIFM**.



### Synthetic Web Service Frontend

Eden required only **three** hints v. **AIFM's** data structure refactoring



### Pandas-like C++ DataFrame

Eden needed **25** hints v. **1,195** LOC changes for running on **AIFM**

## Hinting page faults

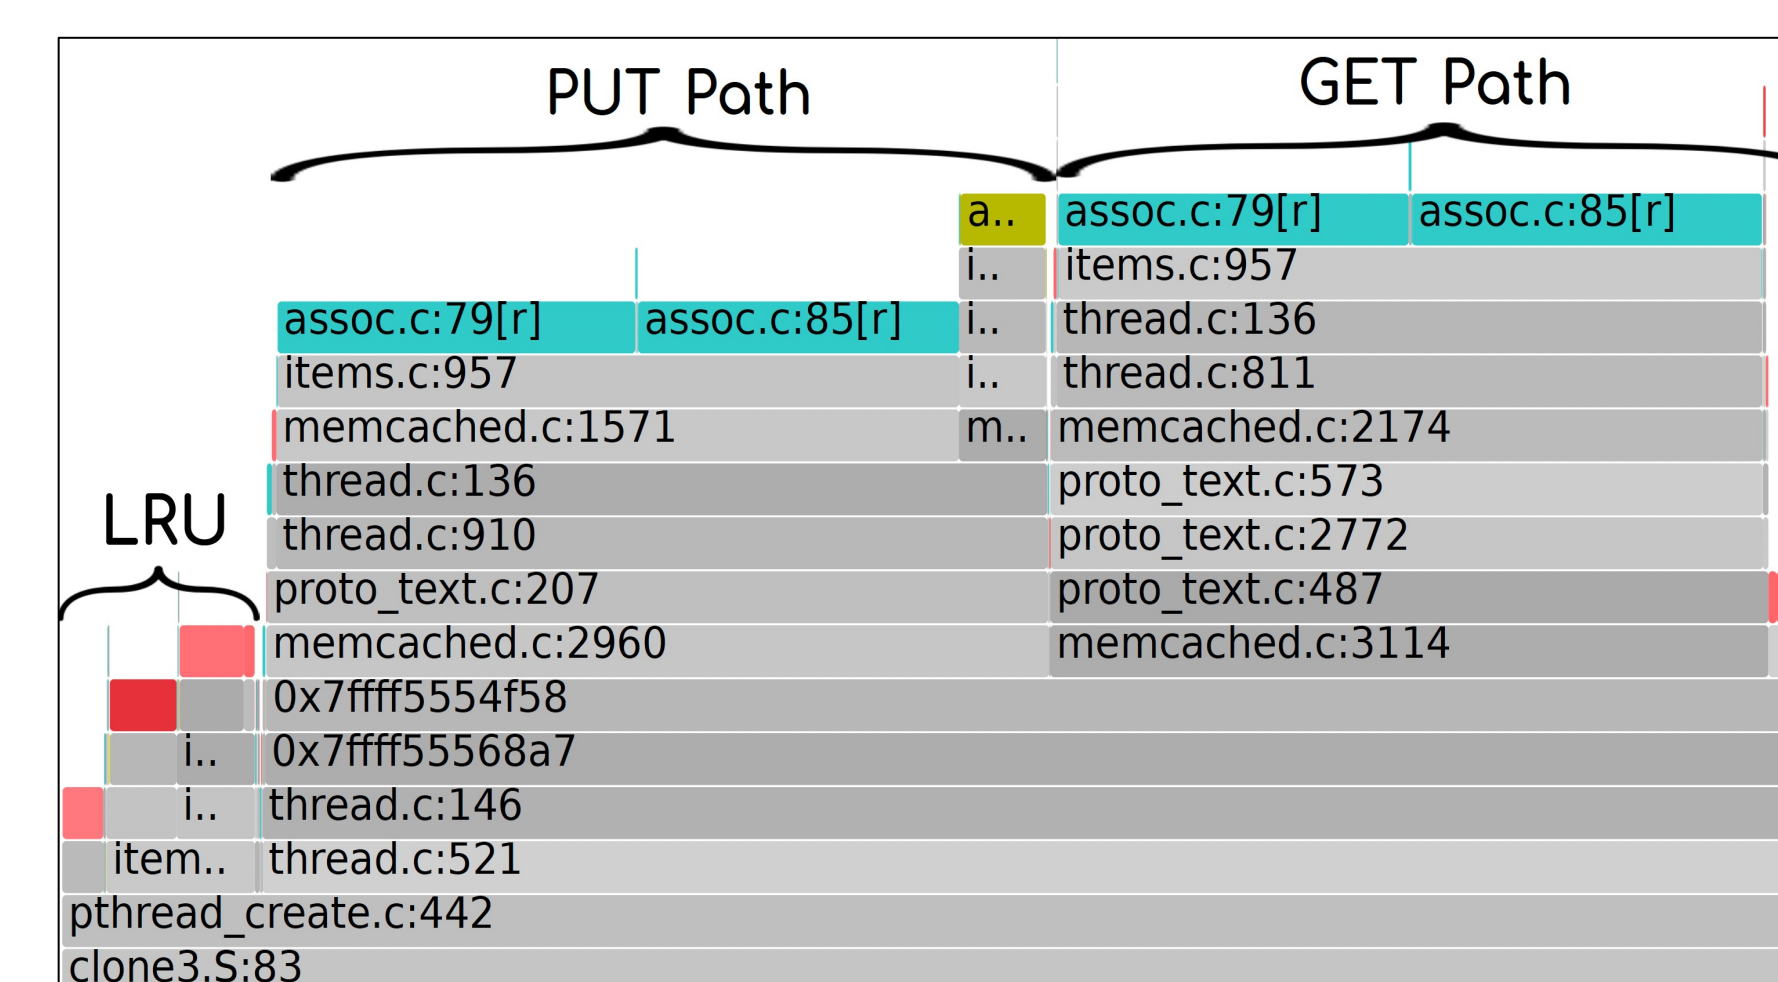
Eden exposes a simple API below for hinting impending page faults

- Eden *requires* at least basic hints
- Extended hints allow application-specific info like **reclaim priority** or **read-ahead**

```
/* basic hints */
hint_page_fault(address, mut = false);
hint_range_fault(address, size, mut = false);

/* extended hints */
hint_(page|range)_fault(
    address, /* faulting address or page */
    size, /* size of the region (for a range fault) */
    mut = false, /* map read-only or writable */
    rdahead = 0, /* positive or negative read-ahead */
    ev_prio = 0, /* eviction priority */
    seq = false, /* sequential access */);
```

Step 1: Run the application binary with our tool `fltrace` to get faulting code locations



Step 2: Add hints at major faulting locations e.g., `assoc.c:85` above

```
82 item *ret = NULL;
83 int depth = 0;
84 while (it) {
85     hint_page_fault(&it->nkey);
86     if ((nkey == it->nkey) && (memcmp(key, ITEM_key(it), nkey) == 0)) {
87         ret = it;
88         break;
89     }
90     it = it->h_next;
```